



RISC-V Memory Protection Unit (MPU)

Dong Du, Xu Lu, Bicheng Yang, Wenhao Li, Yubin Xia, Nick Kossifidis, Joe Xie, Paul Ku, Bill Huffman, Jonathan Behrens, Allen Baum, Robin Zheng, Zeyu Mi, RISC-V TEE Task Group

Version 0.8.0, 09/2021: This document is in the Development state. Assume anything can change. For more information see: <https://riscv.org/spec-state>

Table of Contents

1. Motivation	1
2. Memory Protection Unit(MPU)	2
2.1. Requirements	2
2.2. Memory Protection Unit CSRs	2
2.3. Address Matching	4
2.4. Encoding of Permissions	4
2.5. Priority and Matching Logic	5
2.6. MPU and Paging	5
2.7. Exceptions	6
2.8. Context Switching Optimization	6
3. Summary of Hardware Changes	7
4. Interaction with hypervisor extension	8
4.1. vMPU extension	8
4.2. hgMPU extension	8
5. Interaction with other proposals	10

Chapter 1. Motivation

We propose MPU (RISC-V Memory Protection Unit) to provide isolation when MMU is not available.

RISC-V based processors recently stimulate great interest in the emerging internet of things (IoT). However, as the page-based virtual memory (MMU) is usually not available on IoT devices, it is hard to isolate the S-mode OSes (e.g., RTOS) and user-mode applications. To support secure processing and isolate faults of U-mode software, it is desirable to enable S-mode OS to limit the physical addresses accessible by U-mode software on a hart.

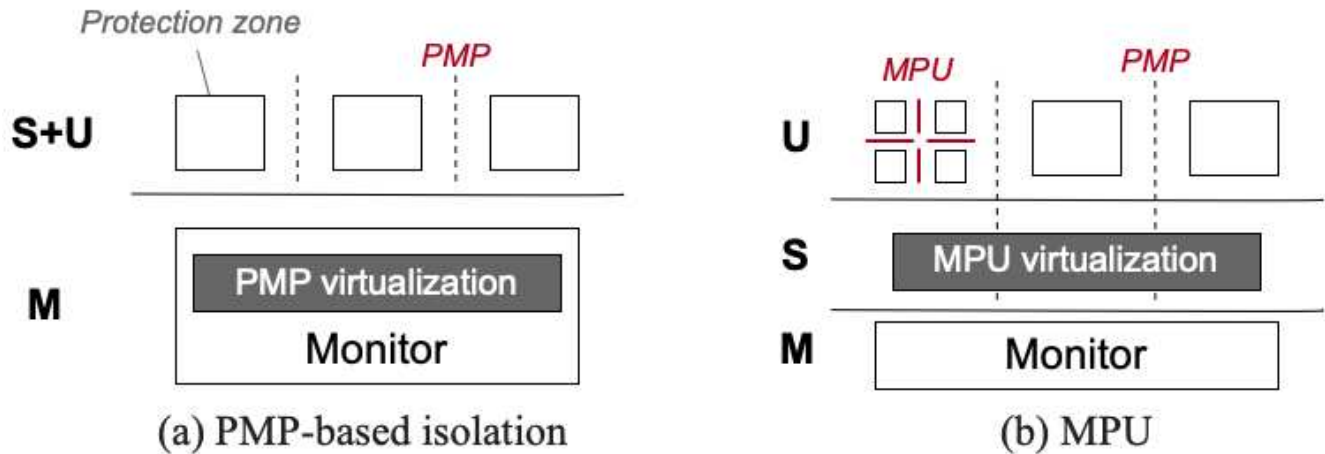


Figure 1. Comparison of PMP and MPU

Chapter 2. Memory Protection Unit(MPU)

An optional RISC-V Memory Protection Unit (MPU) provides per-hart supervisor-mode control registers to allow physical memory access privileges (read, write, execute) to be specified for each physical memory region. The MPU is checked before the PMA checks and PMP checks, the same as MMU.

Like PMP, the granularity of MPU access control settings are platform-specific and within a platform may vary by physical memory region, but the standard MPU encoding should support regions as small as four bytes.

MPU checks will be applied to all accesses for both U mode and S mode, depending on the values in the configuration registers. M-mode accesses are not affected and always pass MPU permission checks. MPU registers can always be modified by M-mode and S-mode software. MPU registers can grant permissions to U-mode, which has none by default, and revoke permissions from S-mode, which has all permissions allowed through PMP/ePMP by default.

2.1. Requirements

1) S mode should be implemented

2.2. Memory Protection Unit CSRs

Like PMP, MPU entries are described by an 8-bit configuration register and one XLEN-bit address register. Some MPU settings additionally use the address register associated with the preceding MPU entry. The number of MPU entries can vary by implementation, and up to 64 MPU entries are supported in standard.



The terms, entry and rule, are similar to ePMP.

The MPU configuration registers are packed into CSRs in the same way as PMP does. For RV32, 16 CSRs, mpucfg0-mpucfg15, hold the configurations mpu0cfg-mpu63cfg for the 64 MPU entries. For RV64, even numbered CSRs (i.e., mpucfg0, mpucfg2, ..., mpucfg14) hold the configurations for the 64 MPU entries; odd numbered CSRs (e.g., mpucfg1) are illegal. Figure 2 and 3 demonstrate the first 16 entries of MPU, the layout of rest entries is similar.

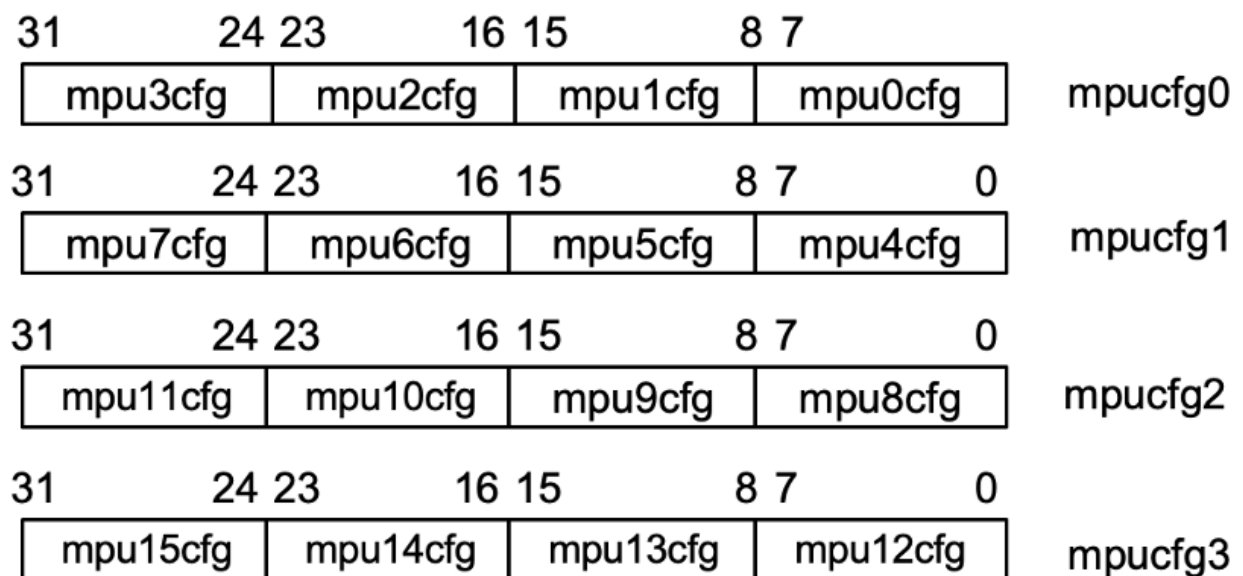


Figure 2. RV32 MPU configuration CSR layout




Figure 3. RV64 MPU configuration CSR layout

The MPU address registers are CSRs named mpuaddr0-mpuaddr63. Each MPU address register encodes bits 33–2 of 34-bit physical address for RV32, as shown in Figure 4. For RV64, each MPU address encodes bits 55–2 of a 56-bit physical address, as shown in Figure 5. Fewer address bits may be implemented for specific reasons, e.g., systems have a smaller physical address space. Implemented address bits must be contiguous and have to go from lower to higher bits.




Figure 4. MPU address register format, RV32



Figure 5. MPU address register format, RV64

The layout of MPU configuration registers is the same as PMP configuration registers, as is shown in Figure 6. The whole register is WARL.

1. The S bit marks a rule as **S-mode-only** when set and **U-mode-only** when unset. The encoding of `mpucfg.RW=01`, and the encoding `mpucfg.SRWX=1111`, now encode a Shared-Region. The rules and encodings for permission are explained in section 2.4, which resembles the encoding of ePMP (except MPU doesn't use locked rules).
2. Bit 5 and 6 are reserved for future use.
3. The A bit will be described in the following sections (2.3).
4. The R/W/X bits control read, write, and instruction execution permissions.




Figure 6. MPU configuration register format

The number of MPU entries: The proposal allows 64 MPU entries, which can provide 64 isolated regions concurrently. To provide more isolation regions, the software in S-mode (usually an OS) can virtualize more isolated regions and schedule them by switching the values in MPU entries.

The reset state: On system reset, the A field of mpu[i]cfg should be zero.

2.3. Address Matching

The A field in an MPU entry's configuration register encodes the address-matching mode of the associated MPU address register. It is the same as PMP/ePMP.

2.4. Encoding of Permissions

MPU has three kinds of rules: **U-mode-only**, **S-mode-only**, and **Shared-Region** rules. The S bit marks a rule as **S-mode-only** when set and **U-mode-only** when unset. The encoding `mpucfg.RW=01` encodes a Shared-Region and `mpucfg.SRWX=1000` is reserved for future standard use.

1. An *S-mode-only* rule is **enforced** on Supervisor mode and **denied** on User mode.
2. A *U-mode-only* rule is **enforced** on User modes and **denied/enforced** on Supervisor mode depending on the value of `sstatus.SUM` bit:
 - If `sstatus.SUM` is set, a U-mode-only rule is enforced yet without code execution permission on Supervisor mode in order to ensure SMEP.
 - If `sstatus.SUM` is unset, a U-mode-only rule is denied on Supervisor mode.
3. A *Shared-Region* rule is enforced on both Supervisor and User modes, with restrictions depending on the `mpucfg.S` and `mpucfg.X` bits:
 - If `mpucfg.S` is not set the region can be used for sharing data between S-mode and U-mode so is not executable. S-mode has RW access to that region and U-mode has read-only access if `mpucfg.X` is not set, or RW access if `mpucfg.X` is set.
 - If `mpucfg.S` is set the region can be used for sharing code between S-mode and U-mode so is not writeable. Both S-mode and U-mode have execute access on the region and S-mode may also have read access if `mpucfg.X` is set.
 - The encoding `mpucfg.SRWX=1111` can be used for sharing data between S-mode and U mode, where both modes only have read-only access to the region.

The encoding and results are shown in the table:

Bits on <i>mpucfg</i> register				Result		
S	R	W	X	S Mode		U Mode
				SUM=0	SUM=1	SUM=0/1
0	0	0	0	Inaccessible region (Access Exception)		
0	0	0	1	Access Exception		Execute-only region
0	0	1	0	Shared data region: Read/write on S mode, read-only on U mode		
0	0	1	1	Shared data region: Read/write for both S and U mode		
0	1	0	0	Access Exception	Read-only region	Read-only region
0	1	0	1	Access Exception	Read-only region	Read/Execute region
0	1	1	0	Access Exception	Read/Write region	Read/Write region
0	1	1	1	Access Exception	Read/Write region	Read/Write/Execute region
1	0	0	0	Reserved		

Bits on <i>mpucfg</i> register				Result	
1	0	0	1	Execute-only region	Access Exception
1	0	1	0	Shared code region: Execute only on both S and U mode	
1	0	1	1	Shared code region: Execute only on U mode, read/execute on S mode	
1	1	0	0	Read-only region	Access Exception
1	1	0	1	Read/Execute region	Access Exception
1	1	1	0	Read/Write region	Access Exception
1	1	1	1	Shared data region: Read only on both S and U mode	

SUM bit: We re-use the `sstatus.SUM` (allow Supervisor User Memory access) bit to modify the privilege with which S-mode loads and stores access physical memory. The semantics of SUM in MPU is consistent with it in paging.

2.5. Priority and Matching Logic

M-mode accesses are always considered to pass MPU checks. If PMP/ePMP is implemented, then accesses succeed only if both PMP/ePMP and MPU permission checks pass.

Like PMP entries, MPU entries are also statically prioritized. The lowest-numbered MPU entry that matches any byte of an access (indicated by an address and the accessed length) determines whether that access is allowed or fails. The matching MPU entry must match all bytes of an access, or the access fails, irrespective of the S, R, W, and X bits.

1. If the privilege mode of the access is M, the access is allowed;
2. If the privilege mode of the access is S and no MPU entry matches, the access is allowed;
3. If the privilege mode of the access is U and no MPU entry matches, but at least one MPU entry is implemented, the access fails;
4. Otherwise, the access is checked according to the permission bits in the matching MPU entry and is allowed only if it satisfies the permission checking with the S, R, W, or X bit corresponding to the access type.

2.6. MPU and Paging

The table below shows which mechanism to use. (Assume both MMU and MPU are implemented.)

satp	Isolation mechanism
<code>satp.mode == Bare</code>	MPU only
<code>satp.mode != Bare</code>	MMU only

We do not allow both MPU and MMU permissions active at the same time now because: (1) It will introduce one more layer to check permission for each memory access. This issue will be more serious for guest OS which may have host MPU and guest MPU. (2) MMU can provide sufficient protection.

That means, MPU is enabled when `satp.mode==Bare` and MPU is implemented.



If page-based virtual memory is not implemented, or when it is disabled, memory accesses check the MPU settings synchronously, so no fence is needed.

2.7. Exceptions

Failed accesses generate an exception. MPU follows the strategy that uses different exception codes for different cases, i.e., load, store/AMO, instruction faults for memory load, memory store/AMO and instruction fetch respectively.

The MPU reuses exception codes of page fault for MPU fault. This is because page fault is typically delegated to S-mode, and so does MPU, so we can benefit from reusing page fault. S-mode software(i.e., OS) can distinguish page fault from MPU fault by checking satp.mode (as mentioned in 2.6, MPU and MMU will not be activated simultaneously). The **MPU is proposing to rename page fault to MPU/MMU fault for clarity**.

Note that a single instruction may generate multiple accesses, which may not be mutually atomic.

Table of renamed exception codes:

Interrupt	Exception Code	Description
0	12	Instruction MPU/MMU fault
0	13	Load MPU/MMU fault
0	15	Store/AMO MPU/MMU fault



You can refer to the Table 3.6 in riscv-privileged spec.

Delegation: Unlike PMP which uses access faults for violations, MPU uses MPU/MMU faults for violations. The benefit of using MPU/MMU faults is that we can delegate the violations caused by MPU to S-mode, while the access violations caused by PMP can still be handled by machine mode.

2.8. Context Switching Optimization

With MPU, each context switch requires the OS to store 64 address registers and 8 configuration registers (RV64), which is costly and unnecessary. So the MPU is proposing an optimization to minimize the overhead caused by context switching.

We add two CSRs called ***mpuswitch0*** and ***mpuswitch1***, which are XLEN-bit read/write registers, formatted as shown in Figure 7. For RV64, only ***mpuswitch0*** is used. Each bit of this register holds on/off status of the corresponding MPU entry respectively. During context switch, the OS can simply store and restore mpuswitch as part of the context. An MPU entry is activated only when both corresponding bits in mpuswitch and A field of mpucfg are set. (i.e., mpuswitch[i] & mpu[i]cfg.A)



Figure 7. MPU domain switch register format (RV64)

Chapter 3. Summary of Hardware Changes

Item	Changes
CSRs for MPU address	64 new CSRs
CSRs for MPU configuration	16 new CSRs for RV32 and 8 for RV64
CSR for Domain switch	2 new CSRs for RV32 and 1 for RV64
Renamed exception code	<i>Instruction page fault</i> renamed to <i>Instruction MPU/MMU fault</i> <i>Load page fault</i> renamed to <i>Load MPU/MMU fault</i> <i>Store/AMO page fault</i> renamed to <i>Store/AMO MPU/MMU fault</i>

Chapter 4. Interaction with hypervisor extension

To support both MPU and hypervisor extension, there are some further changes.

4.1. vMPU extension

This extension describes how MPU is used in a guest VM.

1. A set of vMPU CSRs for the VS-mode are required, including 64 vMPU address registers and 16 configuration registers. When $V=1$, vMPU CSR substitutes for the usual MPU CSR, so instructions that normally read or modify MPU CSR actually access vMPU CSR instead. This is consistent with the paging in VS-mode (i.e., vsatp).
2. For HLV, HLVX, and HSV instructions, the hardware should perform vMPU checking before G-stage address translation (or hgMPU protection when hgatp in BARE mode).
3. The vMPU checking is performed in the guest physical addresses, before G-stage address translation (or hgMPU protection when hgatp in BARE mode).

4.2. hgMPU extension

This extension describes how MPU is used for protecting a hypervisor from guests (only enabled when hgatp is set to BARE mode).

1. When hgMPU is enabled, all guest memory accesses will be checked by hgMPU; while hypervisor (in HS mode) and HU mode applications will not be affected.
2. A set of hgMPU CSRs for the HS-mode are required, including 64 hgMPUaddr address registers and 16 hgMPUcfcg configuration registers. When $V=1$, and hgatp.MODE=Bare, hgMPU is used to provide isolation between hypervisor and guest VMs.
3. XLEN-bit read/write hgmpuswitch0 and hgmpuswitch1 CSRs are also provided in hgMPU, which are identical to mpuswitch0 and mpuswitch1 shown in Figure 7. Only hgmpuswitch0 is used for RV64. During context switch, the hypervisor can simply store and restore hgmpuswitch (we use hgmpuswitch to represent either hgmpuswitch0 or hgmpuswitch1) as part of the context. An hgMPU entry is activated only when both corresponding bits in hgmpuswitch and A field of hgmpuicfcg are set. (i.e., $\text{hgmpuswitch}[i] \ \& \ \text{hgmpuicfcg.A}$)
4. The hgMPU checking is performed after the guest address translation (or vMPU checking), before PMP checking.

As hgMPU does not apply on hypervisor, the encodings of configuration registers are simplified as the following table.

The encodings of hgmpucfcg are shown in the table:

Bits on <i>hgmpucfcg</i> register				Result
S	R	W	X	V Mode (VS + VU)
0	0	0	0	Inaccessible region (Access Exception)
0	0	0	1	Execute-only region
0	1	0	0	Read-only region
0	1	0	1	Read/Execute region

Bits on <i>hgmpucfg</i> register				Result
0	1	1	0	Read/Write region
0	1	1	1	Read/Write/Execute region
Others				Reserved

Chapter 5. Interaction with other proposals

This section discusses how MPU interacts with other proposals.

RISC-V PMP enhancements: MPU is compatible with ePMP proposal, and uses almost the same encoding as ePMP.

J-extension pointer masking proposal: When both PM and MPU are used, MPU checking should be performed using the actual addresses generated by PM (pointer masking).